

ECE421 Lab3 Report

1.K-means

1.1 Learning K-means

1)

distance_func

```
# Distance function for K-means
def distance_func(X, mu):
    """ Inputs:
        X: is an NxD matrix (N observations and D dimensions)
        mu: is an KxD matrix (K means and D dimensions)

        Output:
        pair_dist: is the squared pairwise distance matrix (NxK)
    """
    X = tf.expand_dims(X,1) #shape becomes (N,1,D)
    square = tf.square((X-mu))
    pair_dist = tf.reduce_sum(square,2) #change shape to (N,K)
    return pair_dist
```

Implement K-means

```
def buildGraph(K,D):
    X = tf.placeholder(tf.float64, shape=(None,D))
    mu = tf.Variable(tf.truncated_normal((K, D), mean=0, stddev=1, dtype=tf.float64), trainable=True)
    loss = tf.reduce_sum(tf.reduce_min(distance_func(X,mu), axis=1))
    optimizer = tf.train.AdamOptimizer(learning_rate=0.1, beta1=0.9, beta2=0.99, epsilon=1e-5).minimize(loss)
    return X,mu,loss,optimizer
```

```
def Train_K_means(dataset,K,epochs,valid_data):
    D = dataset.shape[1]
    X,mu,loss,optimizer = buildGraph(K,D)
    loss_List = []
    best_mu = None
    cluster = None
    valid_loss = None
    with tf.Session() as sess:
        sess.run(tf.global_variables_initializer())
        for i in range(epochs):
            best_mu,opt = sess.run([mu,optimizer], feed_dict={X: dataset})
            Loss = sess.run(loss, feed_dict={X: dataset})
            loss_List.append(Loss/dataset.shape[0])

        # end of traning find cluster
        # Returns the index with the smallest value across axes of a tensor
        cluster = sess.run(tf.argmax(distance_func(X, best_mu), 1),feed_dict={X:dataset})
        # end of training find validation loss
        valid_loss = sess.run(loss, feed_dict={X: valid_data})
        valid_loss = valid_loss/valid_data.shape[0]
    return loss_List,best_mu,cluster,valid_loss
```


```

def plot_loss(train, valid = None):
    iterations = range(len(train))
    print("train_loss is ", train[-1])
    plt.plot(iterations, train, label = "training")
    if(valid):
        print("valid_loss is ", valid[-1])
        plt.plot(iterations, valid, label = "validation")
    plt.xlabel("Epoch")
    plt.ylabel("Loss")
    plt.legend()
    plt.title('Loss vs epoch')
    plt.show()
    return

def Q1():
    data = load_dataset(True)
    loss_List, best_mu, cluster, valid_loss = Train_K_means(data, 3, 200, data)
    plot_loss(loss_List)
    return

```

train_loss is 0.511094537291807







2)


```
def plot_cluster(K, dataset, cluster, mu, valid_loss):
    legend = []
    for i in range(K):
        class_i = []
        for j in range(len(cluster)):
            if (cluster[j] == i):
                class_i.append(dataset[j, :])
        class_i = np.array(class_i)
        plt.scatter(class_i[:, 0], class_i[:, 1], cmap='Pastel')
        #calculate number of data belongs to cluster
        percentage = str(round(100*np.sum(i==cluster)/len(cluster), 2))
        legend.append("Cluster"+str(i)+":"+percentage+"%")
    plt.legend(legend)
    plt.scatter(mu[:, 0], mu[:, 1], c='black')
    plt.title("%d-Means Clustering, valid_loss = %f" % (K, valid_loss))
    plt.xlabel('X')
    plt.ylabel('Y')
    plt.grid()
    plt.show()
    return
```

```
def Q2():
    K_list = [1, 2, 3, 4, 5]
    for K in K_list:
        data, val_data = load_dataset(True, True)
        loss_List, best_mu, cluster, valid_loss = Train_K_means(data, K, 100, val_data)
        plot_cluster(K, data, cluster, best_mu, valid_loss)
    return
```

Plot







From the figures above,

Number of Cluster - K	Validation Loss
1	3.861
2	0.888
3	0.488
4	0.316
5	0.275

The minimum validation loss achieved here is when $K=5$, validation loss = 0.275

The validation loss decreases as K increase.

However, it doesn't make sense to divide the dataset into too many clusters.

The validation loss decreases drastically as K increase from 1 to 3, the validation loss drops 3.4.

The validation loss decreases much slower when increase K from 3 to 5, the validation loss drops 2.1.

Thus, I think $K=3$ is the optimal number of clusters.

2. Mixtures of Gaussians

2.1 Gaussian cluster mode

1)

$$N(x, \mu k, \delta k) = \frac{1}{\delta k \sqrt{2\pi}} \exp\left(-\frac{1}{2} \left(\frac{x - \mu k}{\delta k}\right)^2\right)$$

$$\begin{aligned} \log(N(x, \mu k, \delta k)) &= -\log(\delta k \sqrt{2\pi}) - \frac{1}{2} \left(\frac{x - \mu k}{\delta k}\right)^2 \\ &= -\frac{1}{2} \log(2\pi \delta k^2) - \frac{1}{2} \left(\frac{x - \mu k}{\delta k}\right)^2 \end{aligned}$$

```
def log_gauss_pdf(X, mu, sigma):
    """ Inputs:
        X: N X D
        mu: K X D
        sigma: K X 1

    Outputs:
        log Gaussian PDF (N X K)
    """
    part1 = -(0.5) * tf.log(2 * np.pi * tf.transpose(sigma)**2)
    pair_dist = distance_func(X, mu)
    part2 = -(0.5) * tf.square(pair_dist) / (tf.transpose(sigma)**2)
    return part1 + part2
```

2)

$$P(z|x) = \frac{P(x|z)P(z)}{\sum_i^k P(x|i)P(i)}$$

$$\log P(z|x) = \log(P(x|z)P(z)) - \log\left(\sum_i^k P(x|i)P(i)\right)$$

$$= \log P(x|z) + P(z) - \log\left(\sum_i^k P(x|i)P(i)\right)$$

$$= \log P(x|z) + P(z) - \log\left(\sum_i^k \exp(\log P(x|i) + \log P(i))\right)$$

```
def log_posterior(log_PDF, log_pi):
    """ Inputs:
        log_PDF: log Gaussian PDF N X K
        log_pi: K X 1

        Outputs
        log_post: N X K
    """
    log_numerator = log_PDF + tf.squeeze(log_pi)
    log_denominator = reduce_logsumexp(log_numerator, keep_dims=True)
    answer = log_numerator - log_denominator
    return answer
```

The reason to use `reduce_logsumexp` since we need to compute

$$\log \left(\sum_i^k \exp(\log P(x|i) + \log P(i)) \right)$$

while `tf.reduce_sum` only compute the sum

2.2 Learning the MoG

1)

```
def buildGraph(K, D):
    #define variables
    X = tf.placeholder(tf.float32, shape=(None, D))
    mu = tf.Variable(tf.random_normal([K, D], stddev = 1))
    # theta is unconstrained parameter, sigma = exp(phi) with [0 - inf]
    theta = tf.Variable(tf.random_normal([K, 1], stddev = 1))
    sigma = tf.exp(theta)

    # phi is unconstrained parameter, acheive constraint for pi
    phi = tf.Variable(tf.random_normal([K, 1], stddev = 1))
    log_pi = logsoftmax(phi)

    log_pdf = log_gauss_pdf(X, mu, sigma)

    #defien loss & optimizer
    loss= - tf.reduce_sum(reduce_logsumexp(log_pdf + tf.squeeze(log_pi), keep_dims=True))
    optimizer = tf.train.AdamOptimizer(learning_rate=0.1, beta1=0.9, beta2=0.99, epsilon=1e-5).minimize(loss)
    return X, mu, sigma, optimizer, loss, log_pdf, log_pi
```

```

def Train_GMM_means(dataset,K,epochs,valid_data):
    D = dataset.shape[1]
    X,mu,sigma,optimizer,loss,log_pdf,log_pi = buildGraph(K,D)
    loss_List = []
    best_mu = None
    best_sigma = None
    cluster = None
    valid_loss = None
    with tf.Session() as sess:
        sess.run(tf.global_variables_initializer())
        for i in range(epochs):
            best_mu,best_sigma,opt = sess.run([mu,sigma,optimizer], feed_dict={X: dataset})
            Loss = sess.run(loss, feed_dict={X: dataset})
            loss_List.append(Loss/dataset.shape[0])

            # end of traning find cluster
            # Returns the index with the smallest value across axes of a tensor
            cluster = sess.run(tf.argmax(log_posterior(log_pdf,log_pi),1),feed_dict={X:dataset})
            # end of training find validation loss
            valid_loss = sess.run(loss, feed_dict={X: valid_data})
            valid_loss = valid_loss/valid_data.shape[0]

    return loss_List,best_mu,best_sigma,cluster,valid_loss

```


```

def Q1():
    data = load_dataset(True)
    loss_List,best_mu,best_sigma,cluster,valid_loss = Train_GMM_means(data,3,200,data)
    plot_loss(loss_List)
    print("best_mu:",best_mu)
    print("best_sigma",best_sigma)
    return

```

Plot Loss:

train_loss is 1.53743203125





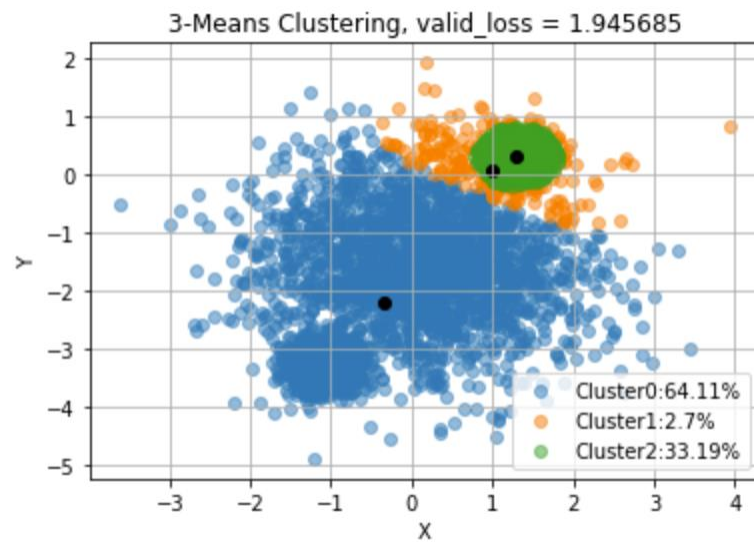
Model Parameter:


Cluster K	Log pi	mu	Sigma
1	-1.0067024	[0.05538136 -1.6027924]	2.9715424
2	-1.1738701	[-1.1191827 -3.316661]	0.07645379
3	-1.1226696	[1.3025047 0.31507367]	0.09755903

2)

```
def Q2():  
    K_list = [1,2,3,4,5]  
    for K in K_list:  
        data, val_data = load_dataset(True, True)  
        loss_List, best_mu, best_sigma, best_pi, cluster, valid_loss = Train_GMM_means(data, K, 100, val_data)  
        plot_cluster(K, data, cluster, best_mu, valid_loss)  
    return
```







From the figures above,

Number of Cluster - K	Validation Loss
1	2.87
2	2.25
3	1.945
4	1.458
5	1.454

The minimum validation loss achieved here is when $K=5$, validation loss = 1.454

The validation loss decreases as K increase.

However, it doesn't make sense to divide the dataset into too many clusters.

The validation loss decreases drastically as K increase from 1 to 3, the validation loss drops 0.93.

The validation loss decreases much slower when increase K from 3 to 5, the validation loss drops 0.39.

Thus, I think $K = 3$ is the optimal number of clusters.

3) Compare K-means and MoG


K	k-means Validation Loss	MoG Validation Loss
5	21.54334444037717	12.697508813381338
10	21.24050911505024	10.834202951545155
15	21.042035154237134	9.654315040879087
20	20.992872165406673	8.715772553817882
30	20.71745446427862	6.316014218609361

The minimum validation loss for both k-means and MoG achieved here is when K=30.

k-means validation loss = 20.717 MoG validation loss = 6.316

The validation loss for both k-means and MoG decreases as K increase.

However, increase in K post more effect on MoG algorithm, since its validation loss decrease more rapidly comparing to the validation loss for k-means as K increase.



From the graph above , we can see that both k-means and Mog validation loss decrease with a steady rate as K increase thus, I think K =30 is the optimal number of clusters.